

# Predicting the Outcome of the 2025 Election in Canada by a Regression Analysis\*

Jingyi Shen

April 2, 2024

ojjojojojljk,mkmk.

## Table of contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Background . . . . .	2
1.2	Hypothesis . . . . .	2
1.3	Terminology . . . . .	2
<b>2</b>	<b>Data</b>	<b>3</b>
2.1	Raw Data . . . . .	3
2.2	Data analysis Tools . . . . .	4
2.3	Variable Description . . . . .	4
2.4	Sample of cleaned Data . . . . .	6
2.5	Measurement . . . . .	11
<b>3</b>	<b>Model</b>	<b>13</b>
3.1	Method . . . . .	13
3.2	Model Set-up . . . . .	14
3.3	Model justification . . . . .	14
3.4	Model Coefficients Interpretation . . . . .	14
3.5	Model Equation . . . . .	14
3.6	Analysis of Residuals . . . . .	15
3.7	Plot of binary logistic regression models . . . . .	15
<b>4</b>	<b>Result</b>	<b>15</b>

---

\*Code and data supporting this analysis are available at: <https://github.com/CSCmaster/Final-Project>

5	Discussion	15
6	Appendix	15
	References	15

# 1 Introduction

## 1.1 Background

The Canadian electoral process, epitomizing democratic principles, empowers citizens to elect parliamentarians who shape Parliament's composition and, consequently, the government formation. In the recent 2021 federal election, Justin Trudeau's Liberal Party secured a majority in the 44th Parliament. Voters' primary concerns include the cost of living, healthcare, climate change, post-pandemic recovery, and immigration, with dissatisfaction towards Trudeau's government prevalent. Trudeau's popularity lags behind UK's Poilievre. The government's ambitious immigration plan aims to admit 500,000 migrants annually by 2025, albeit facing a backlog. The Conservative Party, with a focus on economic prosperity, emphasizes attracting highly skilled immigrants while addressing concerns like living costs and housing issues([Canada, n.d.](#)).

## 1.2 Hypothesis

Based on our analysis of historical trends, governmental policies, and electoral procedures, our hypothesis aligns with the prevailing expectation that the Conservatives will emerge as the likely winners in the 2025 election. **We anticipate that the Conservatives will likely secure the highest number of votes, with the Liberals expected to closely follow in second place.** Our research indicates that the vote-counting process in Canada significantly influences the outcome, as voters do not directly elect a specific prime minister candidate. Initially, all votes are tallied based on first preferences. If a candidate obtains 50% of the vote, they win. If not, the candidate with the fewest votes is eliminated([Canada, n.d.](#)). Then, votes from supporters of the eliminated candidate are redistributed to their second-choice candidate. This iterative redistribution continues until one candidate achieves 50% of the vote. Hence, our objective is to forecast the election probabilities for the Conservative and Liberal parties and compare which is higher([n.d.a](#)).

## 1.3 Terminology

In this paper, I used binary logistic regression models to forecast the likelihood of individuals affiliating with the Conservative or Liberal parties. Binary logistic regression is a statistical

technique tailored for predicting binary outcomes, where the variable under consideration has two possible results. In our study, these outcomes represent the probability of individuals aligning with each of the two parties. Each party’s likelihood will be predicted through separate logistic regression models, with the response variables indicating whether individuals vote for the corresponding party or not. The predictors used in these models include age, province, education, religion, and income before tax. The variable “sex” will be initially included but later removed after AIC model testing, as further elaborated in the model selection section. This regression approach is particularly advantageous when there’s a nonlinear relationship between independent and dependent variables, especially in scenarios with categorical dependent variables like ours. Additionally, we will incorporate post-stratification in our analysis to enhance accuracy. Post-stratification helps refine precision estimates by ensuring that the sample is more representative of the entire population.

## 2 Data

### 2.1 Raw Data

In both datasets, we’ve focused on standardizing variables such as age, sex, province, education, religion, and pre-tax income. Conducting a census typically involves higher expenses and more time since it requires reaching out to every individual in the population. Surveys, however, offer a more cost-effective and quicker alternative, especially for large populations, by collecting data from a select group that represents the larger whole.

**Census Data:** Census data refer to information systematically gathered from the entire population of a defined area at a specific point in time. The census data we have comes from the General Social Survey (GSS), conducted on August 12, 2022. We’ve made necessary adjustments to ensure it matches our survey data. Data collection aims to provide a detailed snapshot of demographic, economic, and social characteristics. The GSS is a nationwide survey program designed to collect data on societal trends, behaviors, and attributes across the entire population, offering insights into the changing dynamics of communities.

**Survey Data:** Survey data, are collected from a sample of the population rather than every individual. This method is often used when it’s impractical or unnecessary to include everyone. Survey data aim to infer the characteristics of the larger population based on the responses of the sample. The survey data, obtained from the Canada Election Study (CES) in 2023, includes responses from over 37,000 participants. The CES focuses specifically on voter behavior, attitudes, and the electoral process within the Canadian context. Unlike the GSS, which aims for a comprehensive overview, the CES targets specific topics of interest to derive insights from a segment of the population.

## 2.2 Data analysis Tools

In this research, the `arrow` (Richardson et al. 2024) package significantly impacts the way large datasets are handled and processed. It provides high-performance reading and writing of data in the Arrow file format, including support for Parquet files, which are highly efficient for storing and querying large datasets. `Tidyr` (Wickham, Vaughan, and Girlich 2024) helps in tidying data, meaning it makes it easier to structure datasets so that they are straightforward to work with. It provides functions to transform data into a tidy format, where each variable forms a column, each observation forms a row, and each type of observational unit forms a table. The `mass` package (Venables and Ripley 2002), It provides a wide range of statistical techniques including linear and nonlinear modeling, statistical tests, time series analysis, classification, and clustering. The package is known for its functions to fit generalized linear models, among many other tools. With `summarytools` (Comtois 2022), users can easily generate frequency tables, descriptive statistics summaries, cross-tabulations, and more. It makes exploratory data analysis more efficient and is particularly useful for preliminary data analysis, ensuring that researchers and analysts can understand their data before moving on to more complex analyses. Data analysis was conducted using the R programming language (R Core Team 2022), renowned for its open-source nature and robust statistical analysis capabilities. Visualization complexities were addressed with the `ggplot2` package (Wickham 2016), which supports the creation of intricate graphics. For data manipulation, `dplyr` (Wickham et al. 2022) was employed, providing a streamlined grammar that simplifies dataset filtering, summarization, and reorganization. Fast and efficient data importing was achieved through the use of the `readr` package (Wickham, Hester, and Bryan 2022). The process of generating this report was seamlessly managed by `knitr` (Xie 2014), facilitating the embedding of R code directly within the text.

## 2.3 Variable Description

### 2.3.1 Survey Data Variable

**Age:** The `cps21_age` field from the survey data denotes participants' ages at the time they completed the General Social Survey. This variable requires no alterations but should be renamed to "age" to ensure consistency with the census data's corresponding variable.

**sex:** In the survey dataset, `cps21_genderid` identifies 9474 respondents as male, 11370 as female, 90 as non-binary, and 34 as other. Given the relatively small number of non-binary and other responses, distribute these respondents between male and female categories based on existing proportions, with females at 54.22% and males at 45.18%.

**province:** The survey data's `cps21_province` includes three additional provinces not present in the census data, creating discrepancies. To align the datasets, remove Northwest Territories, Nunavut, and Yukon from the survey data. Moreover, convert the numeric province labels in the survey data into their corresponding categorical names for analytical compatibility.

**Education:** The survey data's `cps21_education`, featuring twelve education levels, should be consolidated into three categories for analysis. Group No schooling through "Don't know/Prefer not to answer" as "Limited Education"; "Some secondary/ high school" through "Some university as Some Education"; and degrees from "Bachelor's degree" onwards as "Highly Educated".

**Religion:** The `cps21_religion` variable indicates a respondent's religious affiliation. A response of '1' denotes atheism (categorized as 'NO'), whereas any other response signifies a religious affiliation (categorized as 'YES').

**Income\_before\_tax:** Convert the numerical `cps21_income_numbe`, indicating total household income before taxes, into categories for logistic regression analysis. Define "Lower Middle Class to Poor" as incomes < \$50,000, "Middle Class" as incomes between \$50,000 and \$124,999 and "Upper Middle Class to Wealthy" as incomes \$125,000.

**vote\_liberal:** The `cps21_votchoice` variable records intended voting behavior. For logistic regression, recode responses to a binary format, where selecting '1' indicates an intention to vote for the Liberal Party ('1'), and any other selection is recoded as '0'.

**vote\_conservative:** for the Conservative Party, recode `cps21_votchoice` to '1' for respondents who choose '2', indicating a preference for this party, and '0' for all other selections.

### 2.3.2 Census Data Variable

**Age Adjustment:** In the census data, age is presented as a decimal number. For consistency with the survey data, round this value to the nearest whole number.

**Gender Representation:** The gender classification in the census data, listed as either 'Male' or 'Female', aligns perfectly with the survey data, requiring no modifications.

**Provincial Data:** The listing of provinces in the census data, including "Newfoundland" and "Labrador", "Nova Scotia", "Quebec", "Saskatchewan", "Ontario", "Alberta", "British Columbia", "Prince Edward Island", "New Brunswick", and "Manitoba", matches that of the survey data, making no further adjustments necessary.

**Educational Levels:** The census data sorts educational attainment into eight categories, with one category for missing information. Align these with the survey data's education classifications by grouping "Less than high school diploma or its equivalent" and missing data as "Limited Education"; "High school diploma or a high school equivalency certificate", "Trade certificate or diploma", "College, CEGEP, or other non-university certificate or diploma", "University certificate or diploma below the bachelor's level" as "Some Education"; and both "Bachelor's degree", and any university certification above a bachelor's degree as "Highly Educated".

**Religious Affiliation:** The "`religion_has_affiliation`" variable in the census data identifies if an individual has a religious affiliation, is uncertain, claims no religious affiliation, or if

the response is missing. For alignment with the survey data, consolidate into two groups: categorize as YES for those with a religious affiliation, uncertain, or missing responses (assuming missing responses indicate reluctance to disclose religious affiliation), and NO for those without any religious affiliation.

**Income Categories:** The Income\_respondent field in the census dataset uses 7 income brackets to classify respondents' earnings. To harmonize with the survey data's "income\_before\_tax" categories, reclassify these into three broader groups. Designate incomes of "\$125,000 and more" to the "Upper Middle Class to Wealthy" group. Categorize incomes of "\$50,000 to \$74,999", "\$75,000 to \$99,999", and "\$100,000 to \$124,999" as "Middle Class". Assign "Less than \$25,000" and "\$25,000 to \$49,999" incomes to the "Lower Middle Class to Poor" group.

## 2.4 Sample of cleaned Data

### 2.4.1 Data Summary Measures


Table 1: Summary of Numerical Variables Across Census and Survey Data

Variable	Minimum	Q1	Median	Average	Q3	Maximum
Age in Census Data	15	37	54	52.180	67	80
Age in Survey Data	18	36	53	51.300	66	97
Liberal Votes in Survey	0	0	0	0.267	1	1
Conservative Votes in Survey	0	0	0	0.249	0	1


The statistical summary of the census and survey data reveals age distributions with the census data capturing a narrower age range (15 to 80 years, with a median of 54) compared to the broader age span in the survey data (18 to 97 years, with a median of 53). The average ages for the census and survey populations are 52.18 and 51.3 years, respectively. Voting data from the survey indicate that a quarter of the respondents are inclined towards the Liberal Party, as shown by an average of 0.267, while a slightly lower average of 0.249 suggests Conservative support; both parties have 50% of the population not voting for them, with a median of 0. The upper bounds for Liberal and Conservative votes are 1, meaning some respondents indicated a vote for these parties, while the third quartile for Conservative votes indicates less overall support compared to the Liberal Party.

## 2.4.2 Data Visualization

Plot.1 Age Distribution Comparison



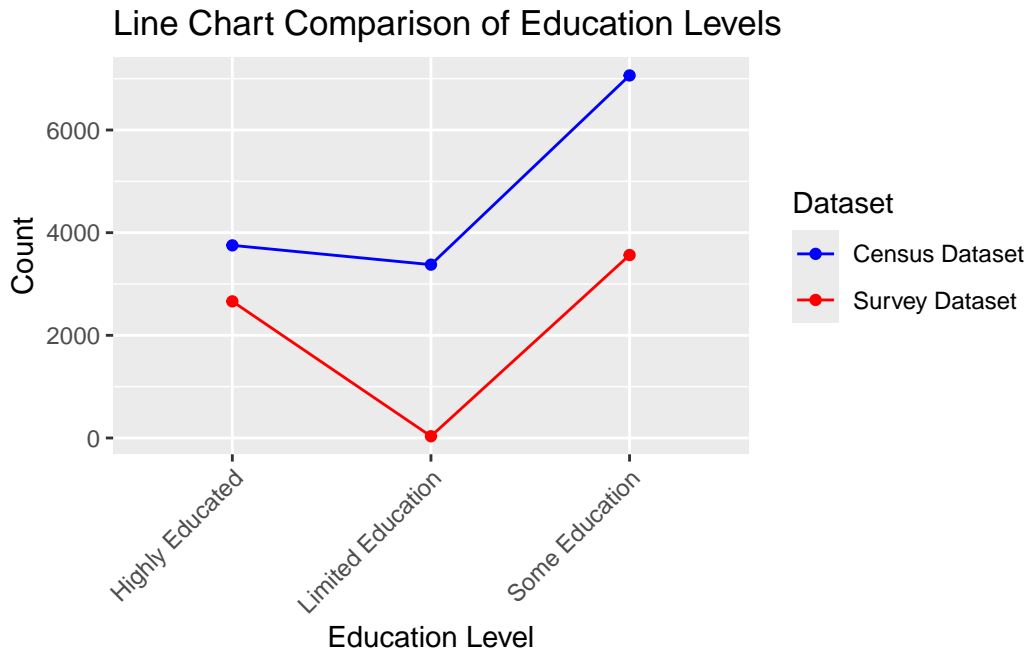
The overlapping histogram illustrates the age distributions from census and survey datasets, with the census data in red displaying a more concentrated and unimodal age distribution peaking between 50 and 60 years and tapering off sharply at 80. In contrast, the survey data, shown in blue, reveals a broader spread of ages with a right-skewed tail extending to 97 years, indicating the inclusion of older respondents. Both distributions share a common median age range, highlighted by the combined purple hue in the overlap, but the census data consistently show higher frequencies across most age bins. This visual comparison underscores the differing age profiles captured by the two datasets and a significant density of middle-aged respondents in both populations.



The bar chart depicts the frequency distribution of respondents across different provinces as collected in two datasets: the census and survey data. Each province is represented on the x-axis, and the corresponding count of respondents is displayed on the y-axis. The bars are color-coded, with blue indicating the survey dataset and red representing the census dataset. The bars are placed side by side, or ‘dodged’, allowing for a direct visual comparison between the two datasets within each provincial category.

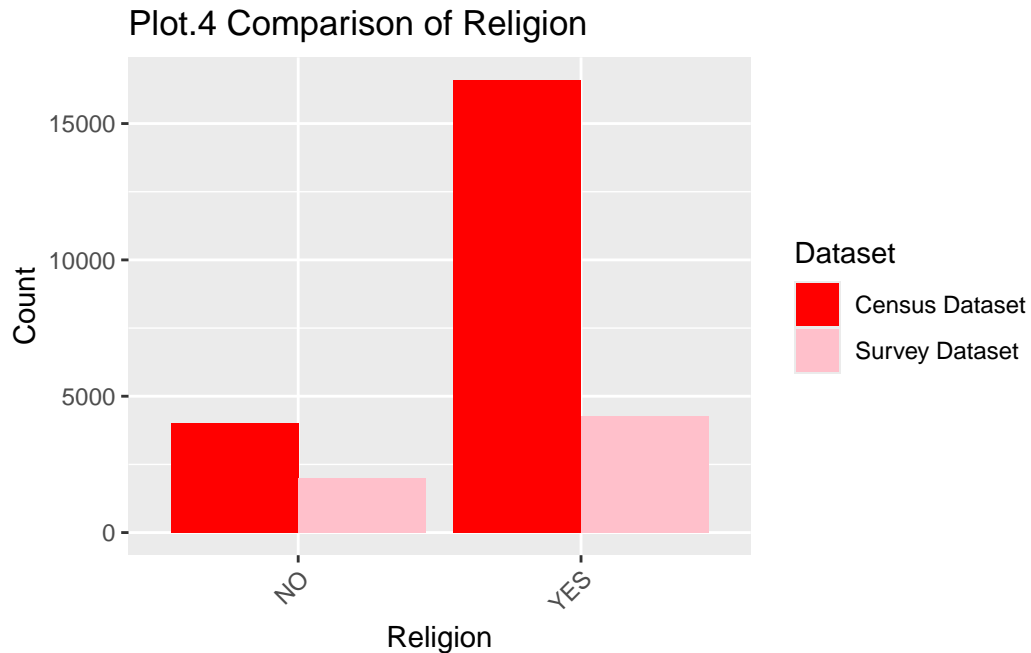
From a statistical perspective, Ontario has the highest frequency in the survey dataset, which substantially exceeds the count in the census dataset, indicating a larger sample or population size recorded in the survey for this province. Conversely, the counts for British Columbia and Alberta show a relatively balanced representation between the two datasets. The chart facilitates a comparative analysis of the datasets, revealing disparities and similarities in the population coverage across provinces, which could be indicative of sampling strategies or the demographic spread within the datasets.






The line chart presented compares the distribution of education levels across two different datasets, one from a survey and another from a census, differentiated by colors—red for the survey and blue for the census. Education levels are divided into three categories along the x-axis: “Highly Educated”, “Limited Education”, and “Some Education”, with the y-axis displaying the count of individuals within each category. A noticeable trend is the inverse relationship between the two datasets from “Highly Educated” to “Limited Education”: the count for the Census dataset declines while that for the Survey dataset rises, suggesting a disparity in the representation of education levels between the two data sources.

This discrepancy highlights potential differences in sampling methods, response rates, or the representativeness of each dataset relative to the underlying population. The most pronounced difference is at the “Limited Education” level, where the Survey dataset shows significantly lower counts compared to the Census dataset, indicating a possible overrepresentation of individuals with limited education in the census or an underrepresentation in the survey. The data suggest that demographic, socioeconomic, or regional factors may influence the education level distribution in each dataset, and further statistical analysis could be warranted to determine the significance and cause of these differences. The differences observed in the chart might also reflect methodological variations in data collection or inherent biases within the sampling frames used.



The bar graph depicts the frequency of individuals categorized by religious affiliation within two datasets: census and survey. Each dataset's frequency is represented by a bar in distinct shades, with a darker red indicating the census dataset and a lighter pink representing the survey dataset. The x-axis bifurcates the data into two groups: those without religious affiliation ('NO') and those with religious affiliation ('YES').

The y-axis provides a count of individuals in each category. Notably, the census dataset shows a substantially higher count of individuals with a religious affiliation compared to those without. This pattern is less pronounced in the survey dataset, where the frequencies of individuals with and without religious affiliation are more comparable. The visual disparity in bar heights, especially within the 'YES' category, suggests a higher prevalence of religious affiliation reported in the census dataset than in the survey dataset.



The bar chart displays the distribution of income categories before tax across two datasets: census and survey. The income categories are defined as ‘Lower Middle Class to Poor’, ‘Middle Class’, and ‘Upper Middle Class to Wealthy’, and are represented on the x-axis. The y-axis measures the count of individuals within each income bracket.

In the ‘Lower Middle Class to Poor’ category, there is a significant count of individuals in the census dataset, shown by a tall bar, which vastly exceeds the count from the survey dataset. For the ‘Middle Class’ bracket, the census dataset again shows a higher frequency, but with a smaller margin compared to the survey dataset. Lastly, in the ‘Upper Middle Class to Wealthy’ category, both datasets depict a lower frequency of individuals than in the other categories, but the survey dataset shows a slightly higher count than the census dataset in this high-income bracket.

The bars for the census dataset are colored in a deep red, while the bars for the survey dataset are in orange, with each pair of bars positioned next to one another for a clear visual comparison. This color differentiation, along with the bar placement, facilitates an immediate comparative analysis of the socioeconomic composition of the populations captured in the two datasets, indicating a notably higher representation of lower income levels within the census data.

## 2.5 Measurement

In this paper, there are two data set, the survey dataset and census dataset. The survey dataset is utilized from Canadian Election Study(CES) while the census dataset is utilized

from General Social Survey(GSS).

Canadian Election Study gathered a sample of 20,968 individuals from across Canada using the Leger Opinion panel, ensuring representation from various regions while maintaining a balance in gender and age demographics within each area. The sample aimed for an equal distribution of 50% males and 50% females, with age groups targeted as follows: 28% between 18-34 years, 33% between 35-54 years, and 39% 55 years and older. The geographical regions included Atlantic Canada, Quebec, Ontario, Western Canada, and the Territories, with quotas set to reflect their population proportions (7% for Atlantic, 23% for Quebec, 38% for Ontario, and 32% for the West) and an open acceptance for any participants from the Territories. In Quebec, the language preference aimed for 80% French-speaking and 20% English-speaking participants, with a goal of 10% French-speaking participants in both the Atlantic region and across Canada. Eligibility for participation required being at least 18 years old and a Canadian citizen or permanent resident(n.d.b).

The General Social Survey focuses on non-institutionalized individuals aged 15 and above residing in Canada's ten provinces. Within each selected household, a random individual is chosen by an application to answer the survey following the initial household roster completion. General Social Survey's questionnaire was developed through rigorous research and broad consultations with data users. The Statistics Canada Questionnaire Design Resource Center (QDRC) conducted qualitative testing in four cities to identify effective questions and those requiring adjustments. Based on the testing, QDRC prepared a comprehensive report with suggestions, which were integrated into the final survey design to enhance clarity and effectiveness. The survey employs a cross-sectional sample survey methodology, utilizing a combined frame of landline and cellular phone numbers from Census data, administrative records, and Statistics Canada's dwelling database. This approach improves coverage of households linked to phone numbers. The sampling strategy is stratified, with stratification occurring at the level of provinces and census metropolitan areas (CMA), and it relies on probability sampling techniques. Information is collected from a single individual aged 15 or older per household, without the use of proxy responses.(n.d.c)

I made several adjustment to the raw dataset to prepare for the analyzed dataset. In the process of cleaning the survey data, several transformative steps were taken to ensure its suitability for analysis. Initially, the age variable was directly adopted as is from `cps21_age`. Gender distribution was then randomized based on predefined proportions for male and female respondents, aiming for a balanced representation. Province information, originally in code form, was converted to textual names for clarity, with any undefined codes being assigned as NA to exclude incomplete records. Educational background was categorized into three broad levels: "Limited Education", "Some Education", and "Highly Educated", based on specific ranges of `cps21_education` codes. Religious affiliation was simplified to a binary "YES"/"NO" based on the `cps21_religion` variable. Income before tax was segmented into three distinct brackets to reflect socioeconomic status. Political preferences were distilled into binary indicators for liberal and conservative vote choices. The dataset was further refined by removing records with missing province information, focusing the analysis on a selected set of variables: age,

sex, province, education, religion, and income before tax. This streamlined dataset was then saved in both CSV and Parquet formats, marking the completion of the data cleaning phase and ensuring a well-structured and analysis-ready dataset.

In cleaning the census data, approach was employed to ensure the data’s accuracy and relevance for analysis. The age variable was rounded to the nearest whole number to standardize age data across the dataset. Education levels were categorized into three distinct groups: “Limited Education” for those with less than a high school diploma or equivalent, including missing data; “Some Education” for individuals with high school diplomas, trade certificates, or college diplomas below the bachelor’s level; and “Highly Educated” for those holding a bachelor’s degree or higher. Religious affiliation was simplified to a binary “YES” for those with any religious affiliation or uncertain/missing responses, and “NO” for those explicitly without religious affiliation. Income levels were stratified into three categories: “Lower Middle Class to Poor” for incomes below \$50,000, “Middle Class” for incomes between \$50,000 and \$124,999, and “Upper Middle Class to Wealthy” for incomes of \$125,000 and above. This process streamlined the dataset by focusing on essential variables for sociodemographic analysis—age, sex, province, education, religion, and income before tax. The cleaned data was then saved in both CSV and Parquet formats to facilitate accessibility and further analysis.

## 3 Model

### 3.1 Method

In forecasting the outcomes of the forthcoming Canadian federal election, statistical techniques like regression analysis, selection of variables, and post-stratification play a pivotal role. Our dataset has some missing variables, but these were deemed negligible in terms of their impact on our predictive model. Initially, the dataset contained 20,921 records, which would have been reduced to 14,544 if entries with missing values were excluded. To maintain the robustness and fullness of our dataset, we opted to retain these incomplete records.

The first step in our modeling approach involved applying logistic regression to estimate the likelihood of a vote being cast for the Liberal or Conservative, based on a set of explanatory variables. Key metrics such as model coefficients and p-values were scrutinized to understand the influence of each predictor. Upon reviewing the p-values, the AIC stepwise selection method was chosen to refine the model further, aiming to lower the AIC score for a more parsimonious fit.

The last stage incorporated post-stratification to gauge the distribution of voter preferences across the identified political factions. This involved segmenting the populace according to our model’s predictors and calculating the voting proportions for each party within these segments.

## 3.2 Model Set-up

## 3.3 Model justification

## 3.4 Model Coefficients Interpretation

## 3.5 Model Equation

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x_{age} + \beta_2 x_{male} + \beta_3 x_{BC} + \beta_4 x_{MB} + \beta_5 x_{NB} + \beta_6 x_{NL} + \beta_7 x_{NS} + \beta_8 x_{ON} + \beta_9 x_{PE} + \beta_{10} x_{QC} + \beta_{11} x_{SK}$$

$p$  is the probability of an individual casting their vote for a certain political party (Liberal or Conservative).

$\beta_0$  also known as the model's intercept, is the base probability of voting in a certain way when all explanatory variables are set to their reference levels.

$\beta_1$  is for every one-unit increase in age, it quantifies the corresponding variation in the log odds of voting preference.

$\beta_2$  is the influence on the log odds of preferring a particular voting option when the person is identified as male, with  $x_{male}$  being assigned a value of 1 for males and 0 for females.

$\beta_3$  to  $\beta_{11}$  is the coefficients detail the adjustments in the log odds of the voting choice associated with residency in different provinces.

$x_{BC}$ , for example, is set to 1 if the person resides in British Columbia, with similar binary indicators for other provinces from  $x_{MB}$  through  $x_{SK}$ .

$\beta_{12}$  captures the shift in log odds of the voting decision when the individual professes a religious belief, with  $x_{religion}$  set to 1 for believers and 0 for non-believers.

$\beta_{13}$  is the modification in the log odds of a particular voting outcome for individuals identified as part of the middle class, where  $x_{mid}$  is 1 for middle-class individuals and 0 otherwise.

$\beta_{13}$  measures the alteration in the log odds of voting behavior for individuals who fall within the upper middle class to wealthy category, with  $x_{upper}$  indicating this socio-economic status by being set to 1, and 0 for others.

### 3.6 Analysis of Residuals

### 3.7 Plot of binary logistic regression models

## 4 Result

## 5 Discussion

## 6 Appendix

## References

- n.d.a. *Elections Step by Step* / Elections Canada's Civic Education. <https://electionsanddemocracy.ca/canadas-elections/canadas-election-process/elections-step-step>.
- . n.d.c. *NORC*. <https://gss.norc.org/>.
- . n.d.b. *Canadian Election Study*. <http://www.ces-eec.ca/>.
- Canada, Elections. n.d. "Past Elections." – *Elections Canada*. <https://www.elections.ca/content.aspx?section=ele&dir=pas&document=index&lang=e>.
- Comtois, Dominic. 2022. *Summarytools: Tools to Quickly and Neatly Summarize Data*. <https://github.com/dcomtois/summarytools>.
- R Core Team. 2022. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Richardson, Neal, Ian Cook, Nic Crane, Dewey Dunnington, Romain François, Jonathan Keane, Dragoş Moldovan-Grünfeld, Jeroen Ooms, Jacob Wujciak-Jens, and Apache Arrow. 2024. *Arrow: Integration to 'Apache' 'Arrow'*. <https://github.com/apache/arrow/>.
- Venables, W. N., and B. D. Ripley. 2002. *Modern Applied Statistics with s*. Fourth. New York: Springer. <https://www.stats.ox.ac.uk/pub/MASS4/>.
- Wickham, Hadley. 2016. *Ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York. <https://ggplot2.tidyverse.org>.
- Wickham, Hadley, Romain François, Lionel Henry, and Kirill Müller. 2022. *Dplyr: A Grammar of Data Manipulation*. <https://CRAN.R-project.org/package=dplyr>.
- Wickham, Hadley, Jim Hester, and Jennifer Bryan. 2022. *Readr: Read Rectangular Text Data*. <https://CRAN.R-project.org/package=readr>.
- Wickham, Hadley, Davis Vaughan, and Maximilian Girlich. 2024. *Tidyr: Tidy Messy Data*. <https://tidyr.tidyverse.org>.
- Xie, Yihui. 2014. "Knitr: A Comprehensive Tool for Reproducible Research in R." In *Implementing Reproducible Computational Research*, edited by Victoria Stodden, Friedrich Leisch, and Roger D. Peng. Chapman; Hall/CRC. <http://www.crcpress.com/product/isbn/9781466561595>.